

Assignment 4

Due at 11:59pm on November 4.

This is an individual assignment. Turn in this assignment as an HTML or PDF file to ELMS. Make sure to include the R Markdown or Quarto file that was used to generate it. Include the GitHub link for the repository containing these files.

GitHub Link: https://github.com/Zackwin73/Assignment_4

In this notebook we will use Google BigQuery, “Google’s fully managed, petabyte scale, low cost analytics data warehouse”. Some instruction on how to connect to Google BigQuery can be found here: <https://db.rstudio.com/databases/big-query/>.

You will need to set up a Google account with a project to be able to use this service. We will be using a public dataset that comes with 1 TB/mo of free processing on Google BigQuery. As long as you do not repeat the work in this notebook constantly, you should be fine with just the free tier.

Go to <https://console.cloud.google.com> and make sure you are logged in a non-university Google account. **This may not work on a university G Suite account because of restrictions on those accounts.** Create a new project by navigating to the dropdown menu at the top (it might say “Select a project”) and selecting “New Project” in the window that pops up. Name it something useful.

After you have initialized a project, paste your project ID into the following chunk.

```
project <- "surv-727-475919"
```

We will connect to a public database, the Chicago crime database, which has data on crime in Chicago.

```
con <- dbConnect(  
  bigrquery::bigrquery(),  
  project = "bigquery-public-data",  
  dataset = "chicago_crime",  
  billing = project
```

```
)  
con
```

```
<BigQueryConnection>  
  Dataset: bigquery-public-data.chicago_crime  
  Billing: surv-727-475919
```

We can look at the available tables in this database using `dbListTables`.

Note: When you run this code, you will be sent to a browser and have to give Google permissions to Tidyverse API Packages. **Make sure you select all to give access or else your code will not run.**

```
dbListTables(con)
```

i Suitable tokens found in the cache, associated with these emails:

```
* 'winogradza@gmail.com'  
  
* 'zackwin@umich.edu'
```

Defaulting to the first email.

! Using an auto-discovered, cached token.

To suppress this message, modify your code or options to clearly consent to the use of a cached token.

See `gargle`'s "Non-interactive auth" vignette for more details:

```
<https://gargle.r-lib.org/articles/non-interactive-auth.html>
```

i The `bigrquery` package is using a cached token for 'winogradza@gmail.com'.

```
[1] "crime"
```

Information on the ‘crime’ table can be found here:

<https://cloud.google.com/bigquery/public-data/chicago-crime-data>

Write a first query that counts the number of rows of the ‘crime’ table in the year 2016. Use code chunks with {sql connection = con} in order to write SQL code within the document.

```
SELECT count(primary_type) AS primary_count, count(*) AS overall_count -- counting non-missing
FROM crime
WHERE year = 2015
LIMIT 10;
```

Table 1: 1 records

primary_count	overall_count
264874	264874

Next, count the number of arrests grouped by `primary_type` in 2016. Note that is a somewhat similar task as above, with some adjustments on which rows should be considered. Sort the results, i.e. list the number of arrests in a descending order.

```
SELECT
    primary_type,
    COUNT(*) AS num_arrests
FROM crime
WHERE year = 2016
    AND arrest = TRUE
GROUP BY primary_type
ORDER BY num_arrests DESC
```

Table 2: Displaying records 1 - 10

primary_type	num_arrests
NARCOTICS	13327
BATTERY	10334
THEFT	6522
CRIMINAL TRESPASS	3724
ASSAULT	3494
OTHER OFFENSE	3416
WEAPONS VIOLATION	2510
CRIMINAL DAMAGE	1669

primary_type	num_arrests
PUBLIC PEACE VIOLATION	1116
MOTOR VEHICLE THEFT	1098

We can also use the `date` for grouping. Count the number of arrests grouped by hour of the day in 2016. You can extract the latter information from `date` via `EXTRACT(HOUR FROM date)`. Which time of the day is associated with the most arrests?

```
SELECT
    EXTRACT(HOUR FROM date) AS hour,
    COUNT(*) AS num_arrests
FROM crime
WHERE year = 2016
    AND arrest = TRUE
GROUP BY hour
ORDER BY hour ASC
```

Table 3: Displaying records 1 - 10

hour	num_arrests
0	2174
1	1556
2	1256
3	980
4	677
5	517
6	781
7	1059
8	1433
9	1714

Focus only on `HOMICIDE` and count the number of arrests for this incident type, grouped by year. List the results in descending order.

```
SELECT
    year,
    COUNT(*) AS num_arrests
FROM crime
WHERE primary_type = 'HOMICIDE'
```

```

    AND arrest = TRUE
GROUP BY year
ORDER BY num_arrests DESC

```

Table 4: Displaying records 1 - 10

year	num_arrests
2001	431
2002	428
2003	386
2020	356
2022	321
2021	296
2004	294
2016	292
2008	288
2006	284

Find out which districts have the highest numbers of arrests in 2015 and 2016. That is, count the number of arrests in 2015 and 2016, grouped by year and district. List the results in descending order.

```

SELECT
    district,
    year,
    COUNT(*) AS arrests
FROM crime
WHERE year IN (2015, 2016)
    AND arrest = TRUE
GROUP BY district, year
ORDER BY SUM(COUNT(*)) OVER (PARTITION BY district) DESC, district, year;

```

Table 5: Displaying records 1 - 10

district	year	arrests
11	2015	8975
11	2016	6578
7	2015	5549
7	2016	3656

district	year	arrests
6	2015	4476
6	2016	3450
15	2015	4514
15	2016	3074
25	2015	4451
25	2016	2953

Lets switch to writing queries from within R via the DBI package. Create a query object that counts the number of arrests grouped by `primary_type` of district 11 in year 2016. The results should be displayed in descending order.

Execute the query.

```
library(DBI)
library(dbplyr)
library(bigrquery)

# Create the SQL query as a string
query <- "
SELECT
    primary_type,
    COUNT(*) AS num_arrests
FROM crime
WHERE year = 2016
    AND district = 11
    AND arrest = TRUE
GROUP BY primary_type
ORDER BY num_arrests DESC
"

arrests_d11_2016 <- dbGetQuery(con, query)

head(arrests_d11_2016)

# A tibble: 6 x 2
  primary_type     num_arrests
  <chr>                <int>
1 NARCOTICS            3634
2 BATTERY                  635
3 PROSTITUTION                 511
```

4 WEAPONS VIOLATION	303
5 OTHER OFFENSE	255
6 ASSAULT	207

Try to write the very same query, now using the `dbplyr` package. For this, you need to first map the `crime` table to a tibble object in R.

```
crime_tbl <- tbl(con, "crime")

arrests_d11_2016_dbplyr <- crime_tbl %>%
  filter(year == 2016,
        district == 11,
        arrest == TRUE) %>%
  group_by(primary_type) %>%
  summarise(num_arrests = n()) %>%
  arrange(desc(num_arrests))

arrests_d11_2016_dbplyr %>% head()
```

```
# Source:      SQL [6 x 2]
# Database:   BigQueryConnection
# Ordered by: desc(num_arrests)
  primary_type      num_arrests
  <chr>                  <int>
1 NARCOTICS            3634
2 BATTERY                635
3 PROSTITUTION           511
4 WEAPONS VIOLATION     303
5 OTHER OFFENSE           255
6 ASSAULT                  207
```

Again, count the number of arrests grouped by `primary_type` of district 11 in year 2016, now using `dplyr` syntax.

```
crime_local <- crime_tbl %>%
  filter(year == 2016, district == 11) %>%
  collect()

arrests_d11_2016_dplyr <- crime_local %>%
  filter(year == 2016,
```

```
district == 11,
arrest == TRUE) %>%
group_by(primary_type) %>%
summarise(num_arrests = n()) %>%
arrange(desc(num_arrests))

head(arrests_d11_2016_dplyr)
```

```
# A tibble: 6 x 2
  primary_type      num_arrests
  <chr>                <int>
1 NARCOTICS            3634
2 BATTERY                 635
3 PROSTITUTION             511
4 WEAPONS VIOLATION        303
5 OTHER OFFENSE              255
6 ASSAULT                  207
```

Count the number of arrests grouped by `primary_type` and `year`, still only for district 11. Arrange the result by `year`.

```
arrests_d11_yearly_query <- crime_tbl %>%
  filter(district == 11, arrest == TRUE) %>%
  group_by(primary_type, year) %>%
  summarise(num_arrests = n()) %>%
  arrange(year)
```

Assign the results of the query above to a local R object.

```
arrests_d11_yearly <- arrests_d11_yearly_query %>% collect()
```

``summarise()` has grouped output by "primary_type". You can override using the ` .groups` argument.`

Confirm that you pulled the data to the local environment by displaying the first ten rows of the saved data set.

```
head(arrests_d11_yearly, 10)
```

```
# A tibble: 10 x 3
# Groups:   primary_type [10]
  primary_type      year num_arrests
  <chr>          <int>     <int>
1 ASSAULT          2001      322
2 OTHER OFFENSE    2001      266
3 KIDNAPPING       2001        4
4 CRIM SEXUAL ASSAULT 2001      17
5 NARCOTICS        2001    7979
6 PUBLIC PEACE VIOLATION 2001      34
7 OFFENSE INVOLVING CHILDREN 2001      44
8 ROBBERY          2001      97
9 GAMBLING          2001      71
10 STALKING         2001       1
```

Close the connection.

```
dbDisconnect(con)
```