

Assignment 4

Due at 11:59pm on November 4.

This is an individual assignment. Turn in this assignment as an HTML or PDF file to ELMS. Make sure to include the R Markdown or Quarto file that was used to generate it. Include the GitHub link for the repository containing these files.

GitHub Link: https://github.com/Zackwin73/Assignment_4

In this notebook we will use Google BigQuery, “Google’s fully managed, petabyte scale, low cost analytics data warehouse”. Some instruction on how to connect to Google BigQuery can be found here: <https://db.rstudio.com/databases/big-query/>.

You will need to set up a Google account with a project to be able to use this service. We will be using a public dataset that comes with 1 TB/mo of free processing on Google BigQuery. As long as you do not repeat the work in this notebook constantly, you should be fine with just the free tier.

Go to <https://console.cloud.google.com> and make sure you are logged in a non-university Google account. **This may not work on a university G Suite account because of restrictions on those accounts.** Create a new project by navigating to the dropdown menu at the top (it might say “Select a project”) and selecting “New Project” in the window that pops up. Name it something useful.

After you have initialized a project, paste your project ID into the following chunk.

```
project <- "surv-727-475919"
```

We will connect to a public database, the Chicago crime database, which has data on crime in Chicago.

```
con <- dbConnect(  
  bigrquery::bigrquery(),  
  project = "bigquery-public-data",  
  dataset = "chicago_crime",  
  billing = project
```

```
)  
con
```

```
<BigQueryConnection>  
  Dataset: bigquery-public-data.chicago_crime  
  Billing: surv-727-475919
```

We can look at the available tables in this database using `dbListTables`.

Note: When you run this code, you will be sent to a browser and have to give Google permissions to Tidyverse API Packages. **Make sure you select all to give access or else your code will not run.**

```
dbListTables(con)
```

! Using an auto-discovered, cached token.

To suppress this message, modify your code or options to clearly consent to the use of a cached token.

See `gargle`'s "Non-interactive auth" vignette for more details:

```
<https://gargle.r-lib.org/articles/non-interactive-auth.html>
```

i The `bigrquery` package is using a cached token for '`winogradza@gmail.com`'.

Auto-refreshing stale OAuth token.

```
[1] "crime"
```

Information on the 'crime' table can be found here:

```
https://cloud.google.com/bigquery/public-data/chicago-crime-data
```

Write a first query that counts the number of rows of the 'crime' table in the year 2016. Use code chunks with `{sql connection = con}` in order to write SQL code within the document.

```
SELECT count(primary_type) AS primary_count, count(*) AS overall_count -- counting non-missing values  
FROM crime  
WHERE year = 2015  
LIMIT 10;
```

Table 1: 1 records

primary_count	overall_count
264873	264873

Next, count the number of arrests grouped by `primary_type` in 2016. Note that is a somewhat similar task as above, with some adjustments on which rows should be considered. Sort the results, i.e. list the number of arrests in a descending order.

```
SELECT
    primary_type,
    COUNT(*) AS num_arrests
FROM crime
WHERE year = 2016
    AND arrest = TRUE
GROUP BY primary_type
ORDER BY num_arrests DESC
```

Table 2: Displaying records 1 - 10

primary_type	num_arrests
NARCOTICS	13327
BATTERY	10334
THEFT	6522
CRIMINAL TRESPASS	3724
ASSAULT	3494
OTHER OFFENSE	3416
WEAPONS VIOLATION	2510
CRIMINAL DAMAGE	1669
PUBLIC PEACE VIOLATION	1116
MOTOR VEHICLE THEFT	1098

We can also use the `date` for grouping. Count the number of arrests grouped by hour of the day in 2016. You can extract the latter information from `date` via `EXTRACT(HOUR FROM date)`. Which time of the day is associated with the most arrests?

```
SELECT
    EXTRACT(HOUR FROM date) AS hour,
    COUNT(*) AS num_arrests
```

```

FROM crime
WHERE year = 2016
    AND arrest = TRUE
GROUP BY hour
ORDER BY hour ASC

```

Table 3: Displaying records 1 - 10

hour	num_arrests
0	2174
1	1556
2	1256
3	980
4	677
5	517
6	781
7	1059
8	1433
9	1714

Focus only on HOMICIDE and count the number of arrests for this incident type, grouped by year. List the results in descending order.

```

SELECT
    year,
    COUNT(*) AS num_arrests
FROM crime
WHERE primary_type = 'HOMICIDE'
    AND arrest = TRUE
GROUP BY year
ORDER BY year DESC

```

Table 4: Displaying records 1 - 10

year	num_arrests
2025	100
2024	213
2023	259
2022	321

year	num_arrests
2021	296
2020	356
2019	200
2018	255
2017	231
2016	292

Find out which districts have the highest numbers of arrests in 2015 and 2016. That is, count the number of arrests in 2015 and 2016, grouped by year and district. List the results in descending order.

```

SELECT
    year,
    district,
    COUNT(*) AS num_arrests
FROM crime
WHERE year IN (2015, 2016)
    AND arrest = TRUE
GROUP BY year, district
ORDER BY num_arrests DESC

```

Table 5: Displaying records 1 - 10

year	district	num_arrests
2015	11	8975
2016	11	6578
2015	7	5549
2015	15	4514
2015	6	4476
2015	25	4451
2015	4	4326
2015	8	4115
2016	7	3656
2015	10	3628

Lets switch to writing queries from within R via the DBI package. Create a query object that counts the number of arrests grouped by `primary_type` of district 11 in year 2016. The results should be displayed in descending order.

Execute the query.

```

library(DBI)
library(dbplyr)
library(bigrquery)

# Create the SQL query as a string
query <- "
  SELECT
    primary_type,
    COUNT(*) AS num_arrests
  FROM crime
  WHERE year = 2016
    AND district = 11
    AND arrest = TRUE
  GROUP BY primary_type
  ORDER BY num_arrests DESC
"

# Run the query and save results
arrests_d11_2016 <- dbGetQuery(con, query)

# Display top rows
head(arrests_d11_2016)

```

	primary_type	num_arrests
	<chr>	<int>
1	NARCOTICS	3634
2	BATTERY	635
3	PROSTITUTION	511
4	WEAPONS VIOLATION	303
5	OTHER OFFENSE	255
6	ASSAULT	207

Try to write the very same query, now using the `dbplyr` package. For this, you need to first map the `crime` table to a tibble object in R.

```

# Map the 'crime' table from BigQuery to a lazy tibble
crime_tbl <-tbl(con, "crime")

# Use dplyr/dbplyr syntax to build the same query
arrests_d11_2016_dbplyr <- crime_tbl %>%

```

```

filter(year == 2016,
       district == 11,
       arrest == TRUE) %>%
group_by(primary_type) %>%
summarise(num_arrests = n()) %>%
arrange(desc(num_arrests))

# Display the first few rows
arrests_d11_2016_dbplyr %>% head()

```

primary_type	num_arrests
NARCOTICS	3634
BATTERY	635
PROSTITUTION	511
WEAPONS VIOLATION	303
OTHER OFFENSE	255
ASSAULT	207

Again, count the number of arrests grouped by `primary_type` of district 11 in year 2016, now using `dplyr` syntax.

```

library(dplyr)

crime_local <- crime_tbl %>%
  filter(year == 2016, district == 11) %>%
  collect()

arrests_d11_2016_dplyr <- crime_local %>%
  filter(year == 2016,
         district == 11,
         arrest == TRUE) %>%
  group_by(primary_type) %>%
  summarise(num_arrests = n()) %>%
  arrange(desc(num_arrests))

head(arrests_d11_2016_dplyr)

```

```
# A tibble: 6 x 2
  primary_type      num_arrests
  <chr>                  <int>
1 NARCOTICS            3634
2 BATTERY                635
3 PROSTITUTION           511
4 WEAPONS VIOLATION     303
5 OTHER OFFENSE          255
6 ASSAULT                 207
```

Count the number of arrests grouped by `primary_type` and `year`, still only for district 11. Arrange the result by `year`.

```
arrests_d11_yearly_query <- crime_tbl %>%
  filter(district == 11, arrest == TRUE) %>%
  group_by(primary_type, year) %>%
  summarise(num_arrests = n()) %>%
  arrange(year)
```

Assign the results of the query above to a local R object.

```
arrests_d11_yearly <- arrests_d11_yearly_query %>% collect()
```

``summarise()`` has grouped output by "primary_type". You can override using the `.groups` argument.

Confirm that you pulled the data to the local environment by displaying the first ten rows of the saved data set.

```
# Display the first ten rows of the locally saved data set
head(arrests_d11_yearly, 10)
```

```
# A tibble: 10 x 3
# Groups:   primary_type [10]
  primary_type              year num_arrests
  <chr>                  <int>      <int>
1 CRIMINAL DAMAGE          2001        163
2 DECEPTIVE PRACTICE        2001         84
3 INTERFERENCE WITH PUBLIC OFFICER 2001        14
4 PROSTITUTION               2001       424
```

5 THEFT	2001	419
6 WEAPONS VIOLATION	2001	236
7 ASSAULT	2001	322
8 STALKING	2001	1
9 NARCOTICS	2001	7979
10 LIQUOR LAW VIOLATION	2001	49

Close the connection.

```
dbDisconnect(con)
```